

**Build your own CV (as in Computer Vision) with  
Keras and Tensorflow**



## Who Am I?

Systems Engineer

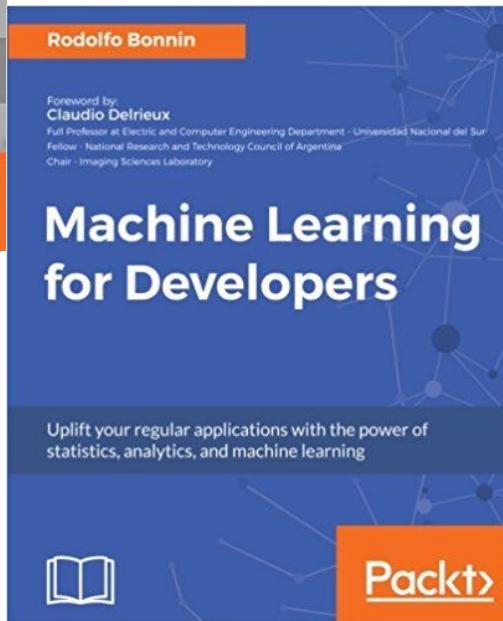
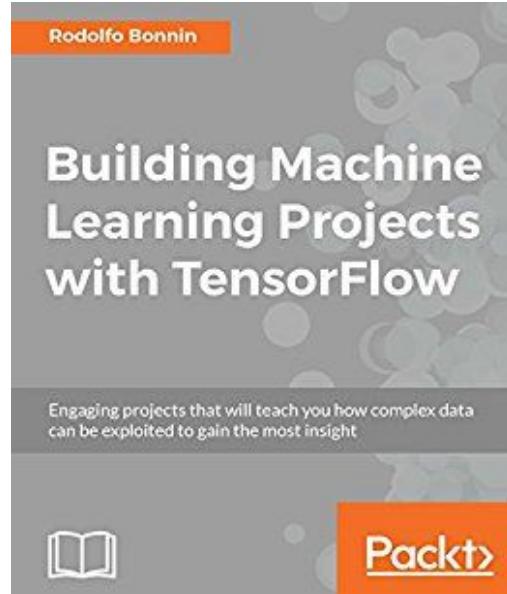
Ph.D Student, UTN

Working on Deep Learning since 2008

Ex ML @ Mercadolibre

CV Specialist @ Machinalis

Author of “Building Machine Learning Projects with Tensorflow” and “Machine Learning for Developers”





# About us

Machinalis

Introduction

Image Classification

Image Detection

Image Segmentation

The Future



## Introduction

## ➤ Evolution of architectures for Computer Vision

1940's

1950's

1960's

1970's

1980's

1990's

2000's

2010's

**Perceptrons**

**First Winter (Minsky & Papert)**

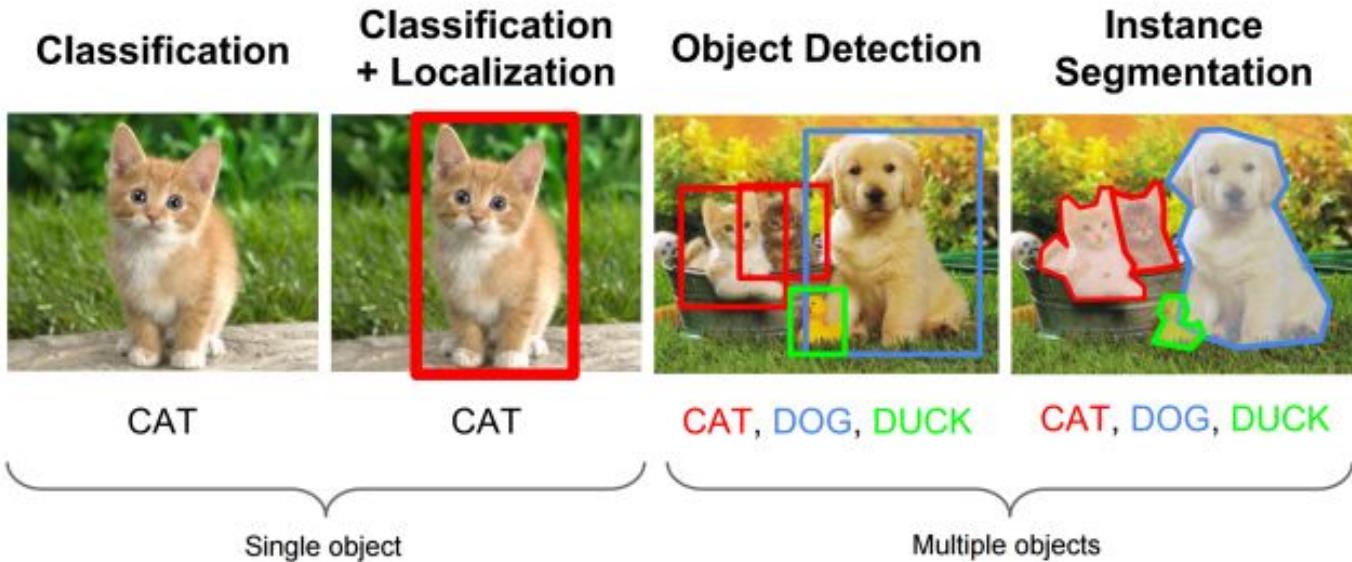
**MLP's**

**First RNN, first Convolutional layers**

**Second winter, SVMs taking place**

**Deep Architectures**

## ➤ Types of tasks for Computer Vision



- **Image classification** is the prediction of the presence/absence of an instance of class in a test image.
- **Object detection** is the prediction of the bounding box and label of each object from the twenty target classes in a test image.
- **Semantic instance segmentation** is the assignment of one of the several labels to every pixel in a test image.



➤ Keras: An easy way to  
build neural models

# > Keras

Keras is a high-level neural networks API, written in Python and capable of running on top of TensorFlow or CNTK.

- Allows for easy and fast prototyping
- Supports both convolutional networks and recurrent networks, as well as combinations of the two) and others via complementary projects.
- Runs seamlessly on CPU and GPU.

# ➤ Model compilation

A basic Keras model needs

- An **optimizer**. This could be the string identifier of an existing optimizer (such as rmsprop or adagrad), or an instance of the Optimizer class.
- A **loss function**. This is the objective that the model will try to minimize. It can be the string identifier of an existing loss function (such as categorical\_crossentropy or mse), or it can be an objective function.
- A list of **metrics**.

```
model.compile(loss='mean_squared_error',
               optimizer='sgd',
               metrics=[metrics.mae, metrics.categorical_accuracy])
```

## ➤ Keras Model Training

Keras models are trained on Numpy arrays of input data and labels.

For training a model, you will typically use the *fit* function.

# > Main Layer Types

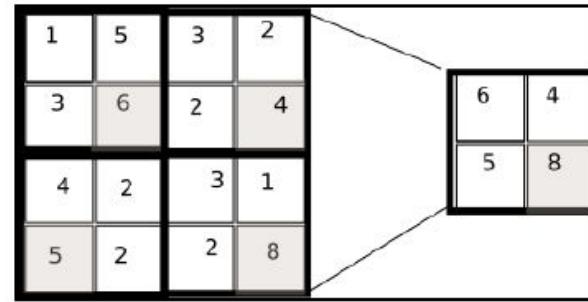
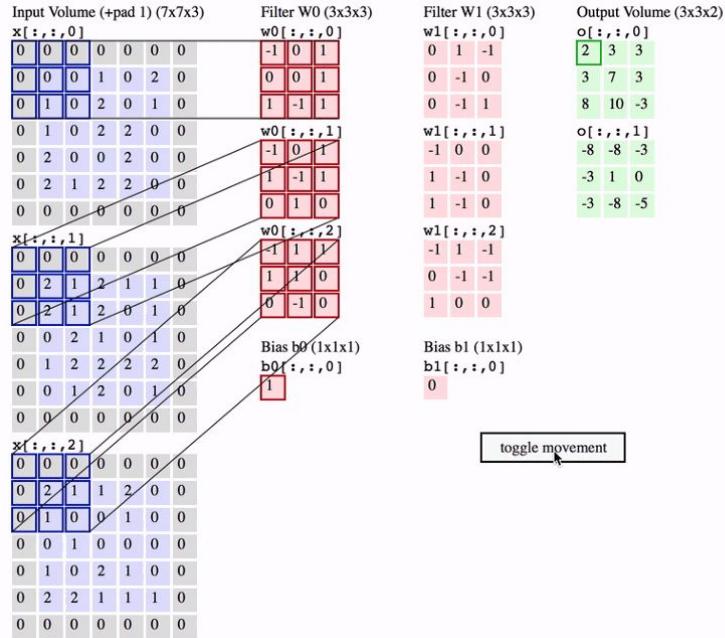
Dense

Convolutional: Conv1D, Conv2D ,Conv3D

Pooling:

Recurrent: RNN, GRU, LSTM

# > Main special layers types in Computer Vision



Pooling

## Convolutional

# > Losses

mean\_squared\_error

mean\_squared\_error(y\_true, y\_pred)

mean\_absolute\_error(y\_true, y\_pred)

mean\_absolute\_percentage\_error(y\_true, y\_pred)

mean\_squared\_logarithmic\_error(y\_true, y\_pred)

categorical\_crossentropy(y\_true, y\_pred)

# > Metrics

- accuracy
- binary\_accuracy
- binary\_accuracy(y\_true, y\_pred)
- categorical\_accuracy(y\_true, y\_pred)
- top\_k\_categorical\_accuracy(y\_true, y\_pred, k=5)

# > Optimizers

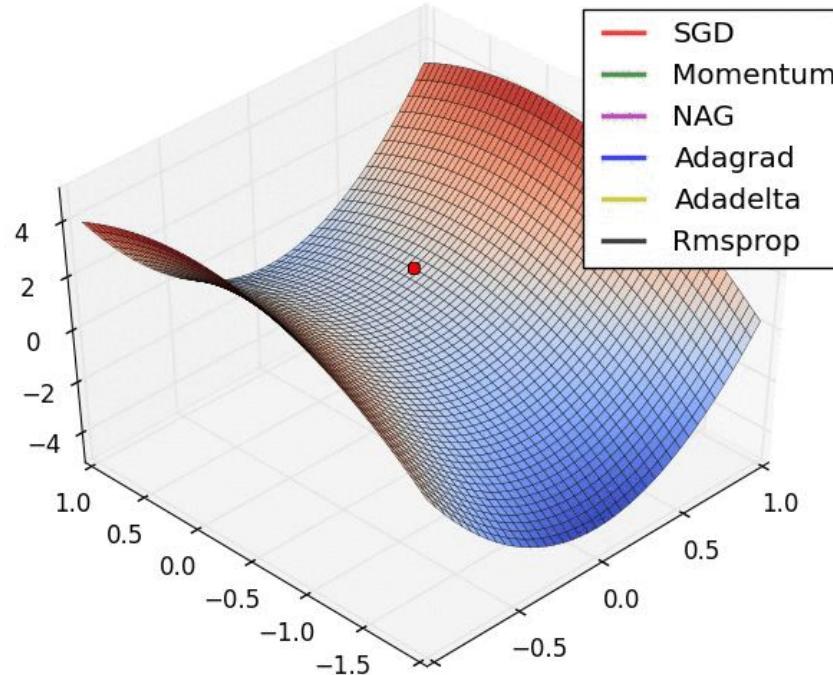
SGD()

RMSprop()

Adagrad()

Adadelta()

Adam()



Source: Alec Radford

# > Activations

softmax

relu

tanh

sigmoid

hard\_sigmoid

linear

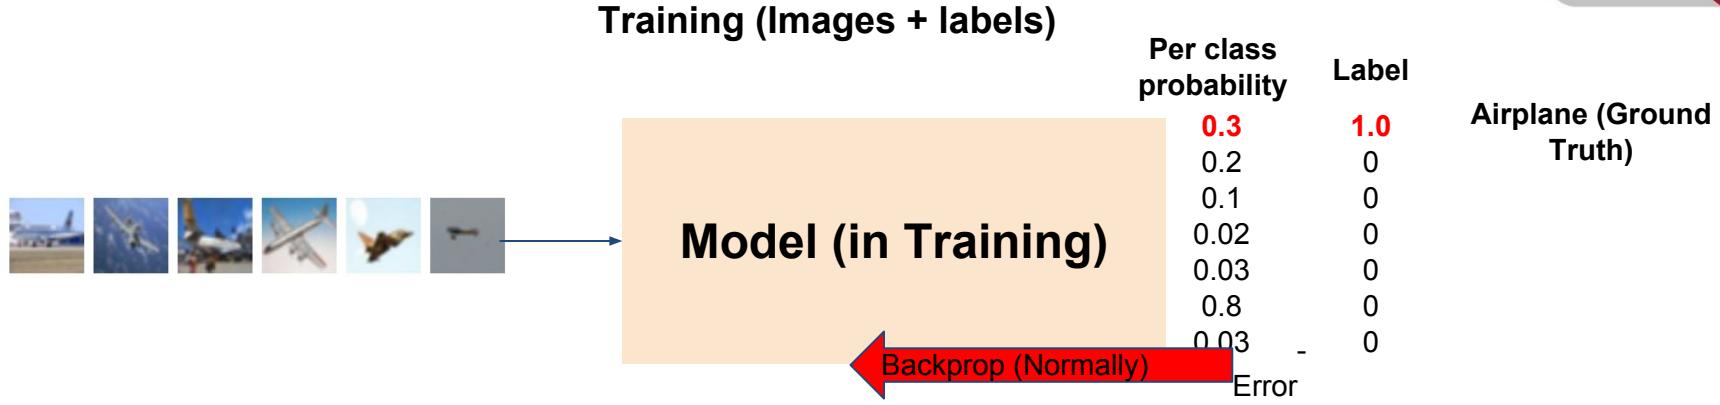
## Example:

```
from keras.layers import Activation, Dense  
model.add(Dense(64))  
model.add(Activation('tanh'))
```



**Image classification:  
History of neural  
models**

## Classification



## From Numbers to clothes

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
3 3 3 3 3 3 3 3 3 3 3 3 3 3 3  
4 4 4 4 4 4 4 4 4 4 4 4 4 4 4  
5 5 5 5 5 5 5 5 5 5 5 5 5 5 5  
6 6 6 6 6 6 6 6 6 6 6 6 6 6 6  
7 7 7 7 7 7 7 7 7 7 7 7 7 7 7  
8 8 8 8 8 8 8 8 8 8 8 8 8 8 8  
9 9 9 9 9 9 9 9 9 9 9 9 9 9 9



## ➤ First notebook: 1\_History\_of\_neural\_models.ipynb

Open in Colaboratory:

[https://colab.research.google.com/github/rnditdev/PyData\\_London\\_2018\\_Computer\\_Vision/blob/master/1\\_History\\_of\\_neural\\_models.ipynb](https://colab.research.google.com/github/rnditdev/PyData_London_2018_Computer_Vision/blob/master/1_History_of_neural_models.ipynb)

### ▼ Fashion Mnist: History of neural models

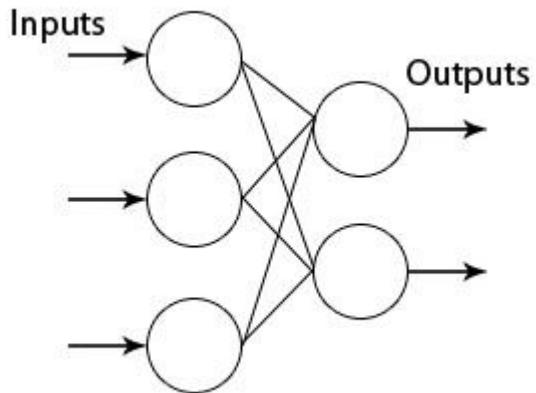
```
▶ from __future__ import print_function
    import keras
    from keras.datasets import fashion_mnist
    from keras.models import Sequential
    from keras.layers import Dense, Dropout, Conv2D, MaxPooling2D, Flatten
    from keras.optimizers import RMSprop
    from keras.optimizers import SGD
    import matplotlib.pyplot as plt
    %matplotlib inline
    batch_size = 128
    num_classes = 10
    epochs = 20

    # the data, shuffled and split between train and test sets
    (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

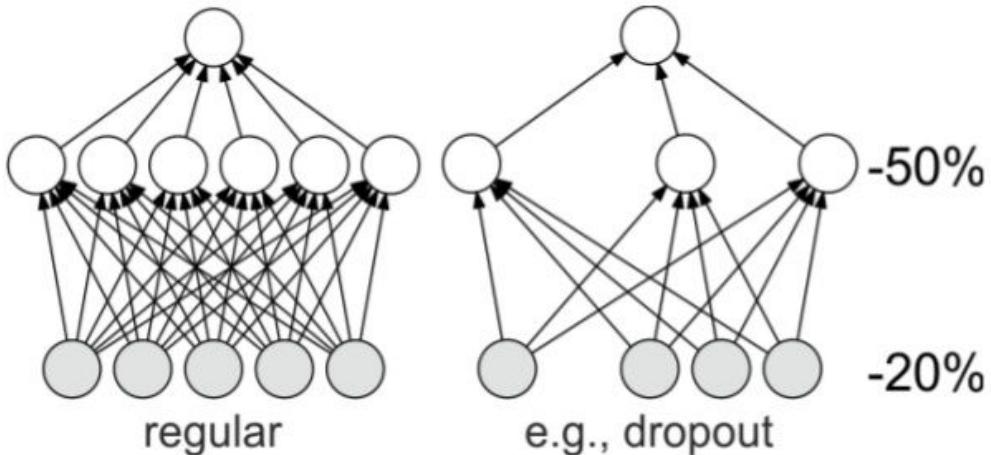
    print(y_train[1000])
    fig= plt.figure()
    plt.imshow(x_train[1000], cmap='gray')

    print(y_train[4000])
    fig= plt.figure()
```

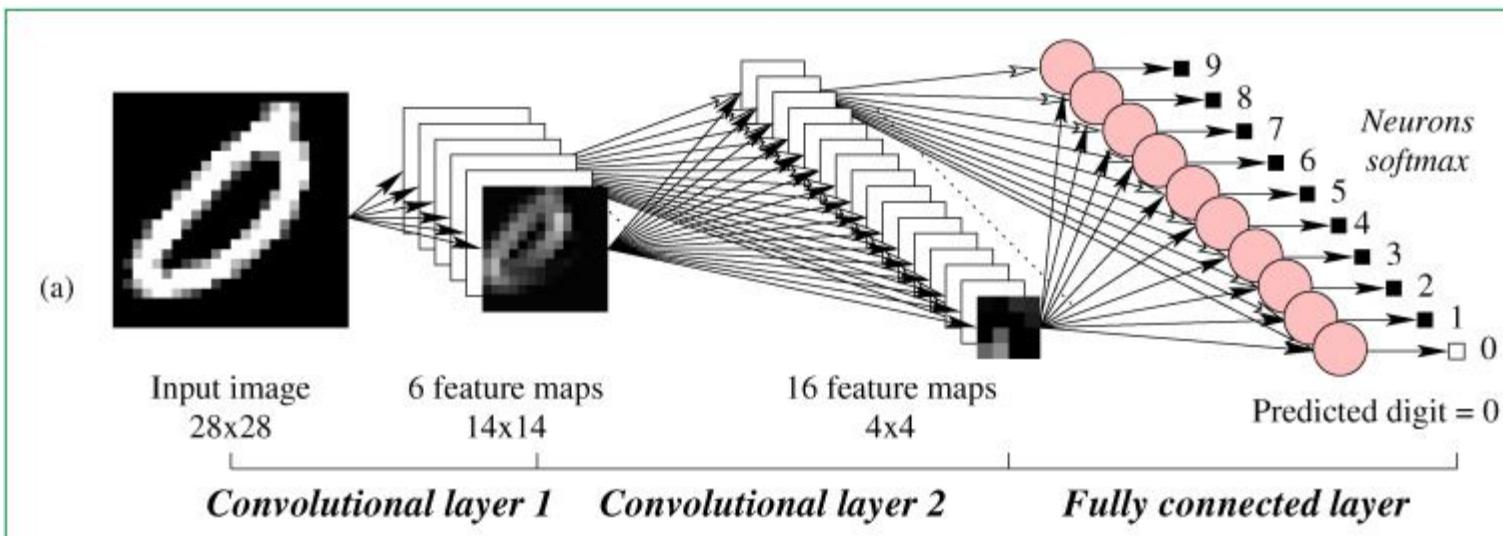
## From Numbers to clothes: Single layer Perceptron



## From Numbers to clothes



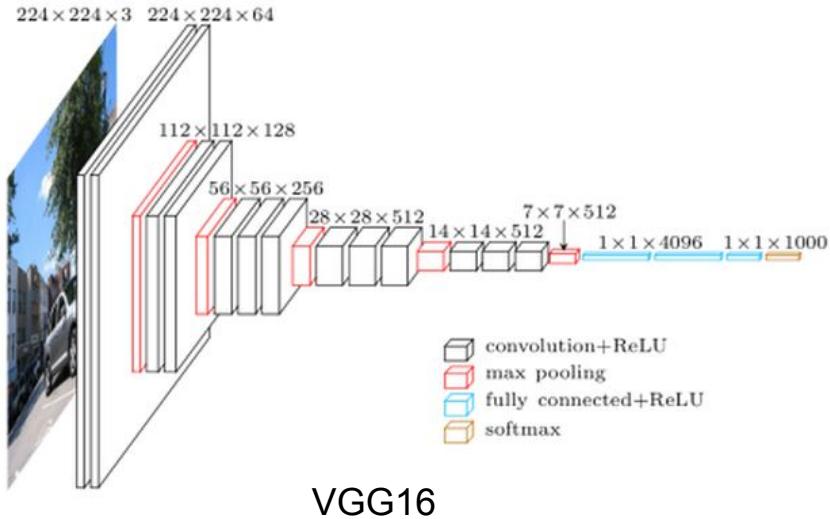
## From Numbers to clothes: Convolutional layer



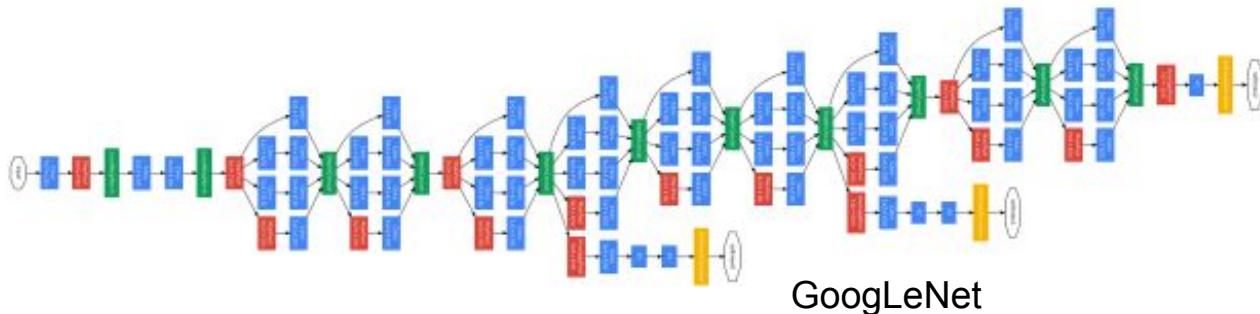


# Deep Learning and Transfer Learning Classification

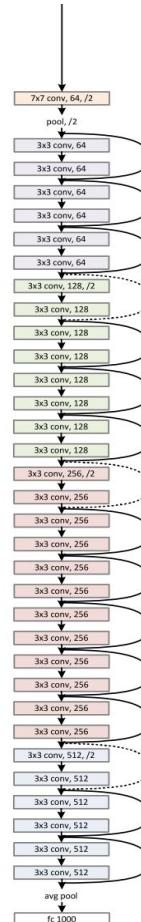
## ➤ Main Deep learning models



VGG16

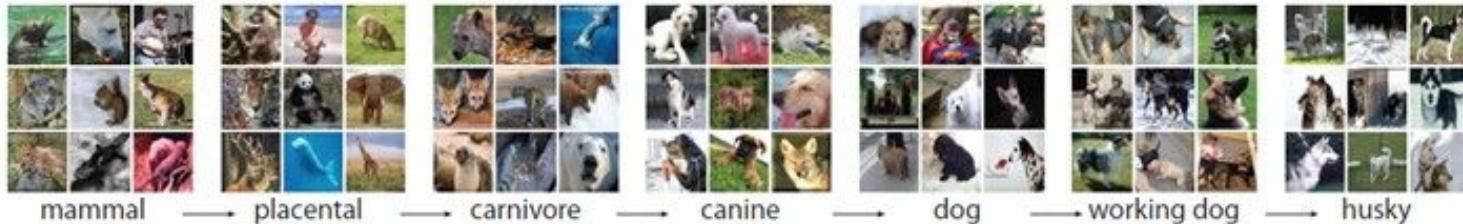


GoogLeNet



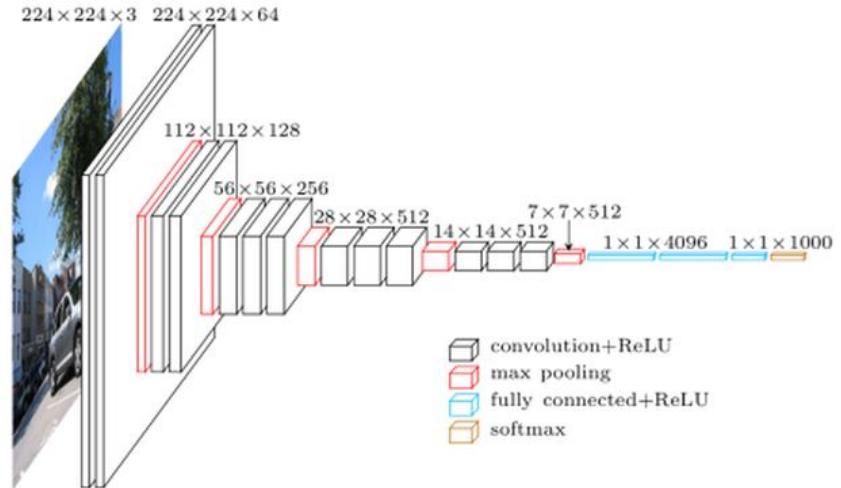
ResNet

# First things: Imagenet



- S: (n) [Eskimo dog](#), [husky](#) (breed of heavy-coated Arctic sled dog)
  - [direct hypernym](#) / [inherited hypernym](#) / [sister term](#)
- S: (n) [working dog](#) (any of several breeds of usually large powerful dogs bred to work as draft animals and guard and guide dogs)
  - S: (n) [dog](#), [domestic dog](#), [Canis familiaris](#) (a member of the genus *Canis* (probably descended from the common wolf) that has been domesticated by man since prehistoric times; occurs in many breeds) "the dog barked all night"
  - S: (n) [canine](#), [canid](#) (any of various fissiped mammals with nonretractile claws and typically long muzzles)
  - S: (n) [carnivore](#) (a terrestrial or aquatic flesh-eating mammal) "terrestrial carnivores have four or five clawed digits on each limb"
  - S: (n) [placental](#), [placental mammal](#), [eutherian](#), [eutherian mammal](#) (mammals having a placenta; all mammals except monotremes and marsupials)
  - S: (n) [mammal](#), [mammalian](#) (any warm-blooded vertebrate having the skin more or less covered with hair; young are born alive except for the small subclass of monotremes and nourished with milk)
  - S: (n) [vertebrate](#), [craniate](#) (animals having a bony or cartilaginous skeleton with a segmented spinal column and a large brain enclosed in a skull or cranium)
  - S: (n) [chordate](#) (any animal of the phylum Chordata having a notochord or spinal column)
    - S: (n) [animal](#), [animate being](#), [beast](#), [brute](#), [creature](#),  [fauna](#) (a living organism characterized by voluntary movement)
    - S: (n) [organism](#), [being](#) (a living thing that has (or can develop) the ability to act or function independently)
    - S: (n) [living thing](#), [animate thing](#) (a living (or once living) entity)
      - S: (n) [whole](#), [unit](#) (an assemblage of parts that is regarded as a single entity) "how big is that part compared to the whole?"; "the team is a unit"
      - S: (n) [object](#), [physical object](#) (a tangible and visible entity; an entity that can cast a shadow) "it was full of rackets, balls and other objects"
      - S: (n) [physical entity](#) (an entity that has physical existence)
      - S: (n) [entity](#) (that which is perceived or known or inferred to have its own distinct existence (living or nonliving))

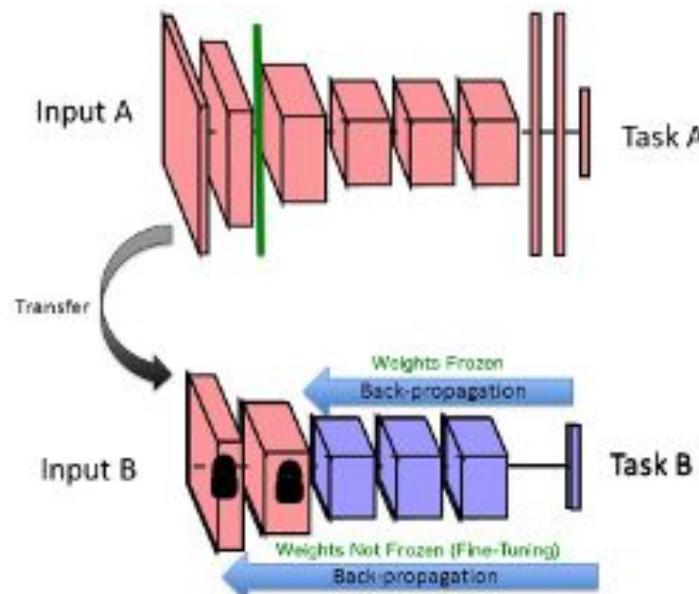
## ➤ Demo time! Exploring CNN layers on a Deep Network with Quiver



## Applying Transfer Learning to build a classifier

Open in Colaboratory:

[https://colab.research.google.com/github/rnditdev/PyData\\_London\\_2018\\_Computer\\_Vision/blob/master/3\\_Transfer%20Learning%20For%20classification.ipynb](https://colab.research.google.com/github/rnditdev/PyData_London_2018_Computer_Vision/blob/master/3_Transfer%20Learning%20For%20classification.ipynb)





➤ Tensorflow &  
Tensorflow Object  
Detection API

## ➤ Tensorflow Object detection API



➤ Object localization  
& Detection

## ➤ The PASCAL Visual Object Classes (2005/2012)

- A publicly available dataset of images and annotations:
  - 2.5 million labeled instances in 328k images, taken from consumer photographs
  - 20 Object types
- An annual competition and workshop, held until 2012
-

# ➤ The PASCAL Visual Object Classes (2005/2012)

- Example annotation:

```
<annotation>
    <folder>GeneratedData_Train</folder>
    <filename>000001.png</filename>
    <path>/my/path/GeneratedData_Train/000001.png</path>
    <source>
        <database>Unknown</database>
    </source>
    <size>
        <width>224</width>
        <height>224</height>
        <depth>3</depth>
    </size>
    <segmented>0</segmented>
    <object>
        <name>21</name>
        <pose>Frontal</pose>
        <truncated>0</truncated>
        <difficult>0</difficult>
        <occluded>0</occluded>
        <bndbox>
            <xmin>82</xmin>
            <xmax>172</xmax>
            <ymin>88</ymin>
            <ymax>146</ymax>
        </bndbox>
    </object>
</annotation>
```

## ➤ Types of detectors

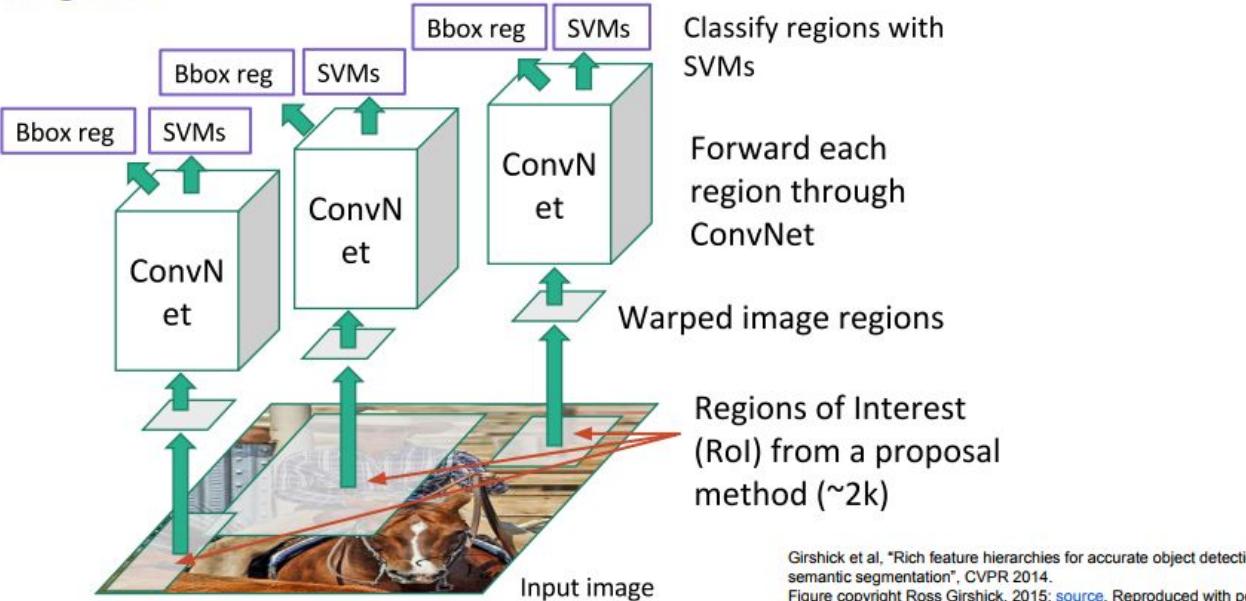
Main types of deep neural  
Detection Architectures

With proposals:  
R-CNN, Faster-CNN

Without Proposals:  
YOLO, SSD

## R-CNN variants

### R-CNN

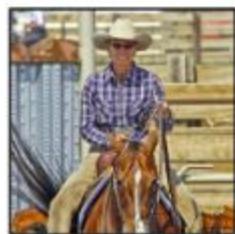


R-CNN creates bounding boxes, or region proposals, using a process called Selective Search which you can read about [here](#). At a high level, Selective Search looks at the image through windows of different sizes, and for each size tries to group together adjacent pixels by texture, color, or intensity to identify objects

<https://blog.athelas.com/a-brief-history-of-cnns-in-image-segmentation-from-r-cnn-to-mask-r-cnn-34ea83205de4>

Thanks CS231

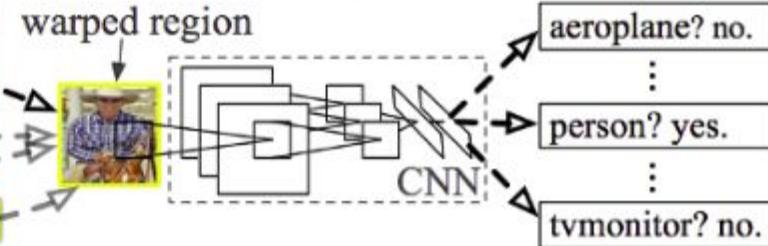
### R-CNN: *Regions with CNN features*



1. Input image



2. Extract region proposals (~2k)



3. Compute CNN features

4. Classify regions

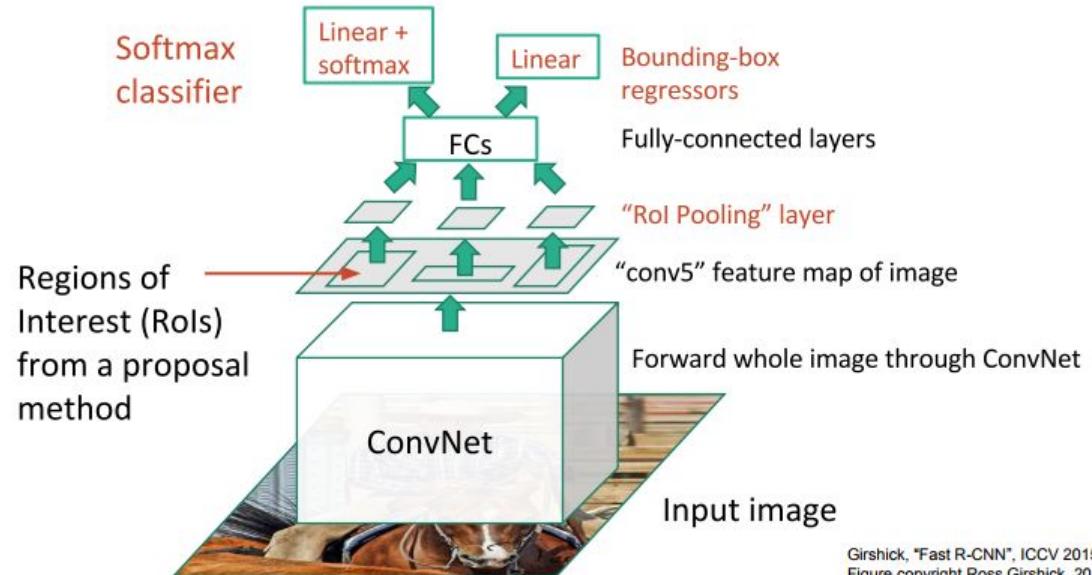
After creating a set of region proposals, R-CNN passes the image through a modified version of AlexNet to determine whether or not it is a valid region. Source: <https://arxiv.org/abs/1311.2524>.

## ➤ R-CNN variants Fst R-CNN

R-CNN works really well, but is slow for a few simple reasons:

- It requires a forward pass of the CNN (AlexNet) for every single region proposal for every single image (that's around 2000 forward passes per image!).
- It has to train three different models separately - the CNN to generate image features, the classifier that predicts the class, and the regression model to tighten the bounding boxes. This makes the pipeline extremely hard to train.

## Fast R-CNN



Girshick, "Fast R-CNN", ICCV 2015.  
Figure copyright Ross Girshick, 2015;

- RoIPool shares the forward pass of a CNN for an image across its subregions.
- One pass of the original image as opposed to ~2000
- Where earlier we had different models to extract image features (CNN), classify (SVM), and tighten bounding boxes (regressor), **Fast R-CNN instead used a single network to compute all three.**

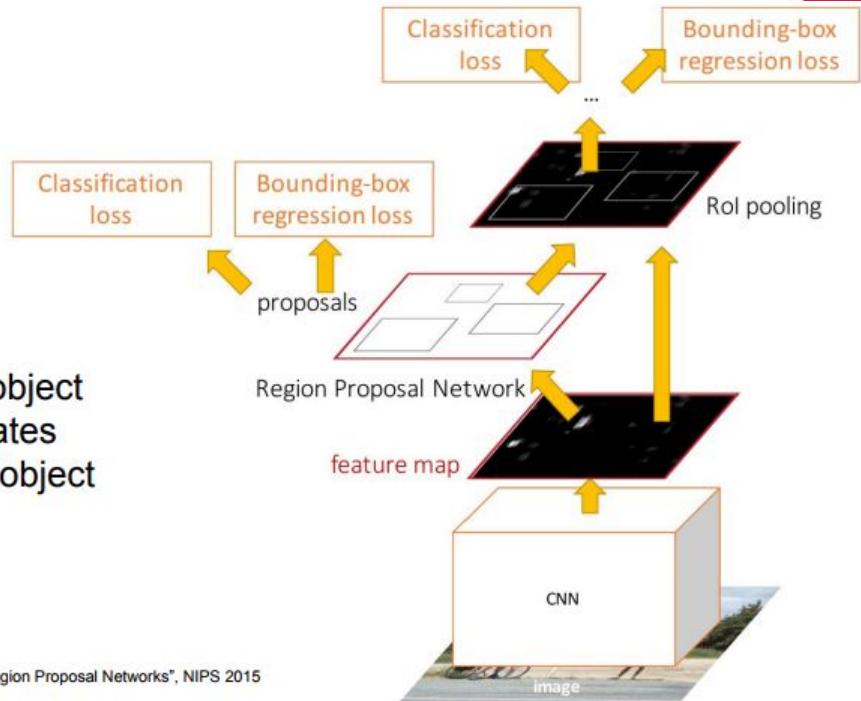
### Faster R-CNN: Make CNN do proposals!

Insert **Region Proposal Network (RPN)** to predict proposals from features

Jointly train with 4 losses:

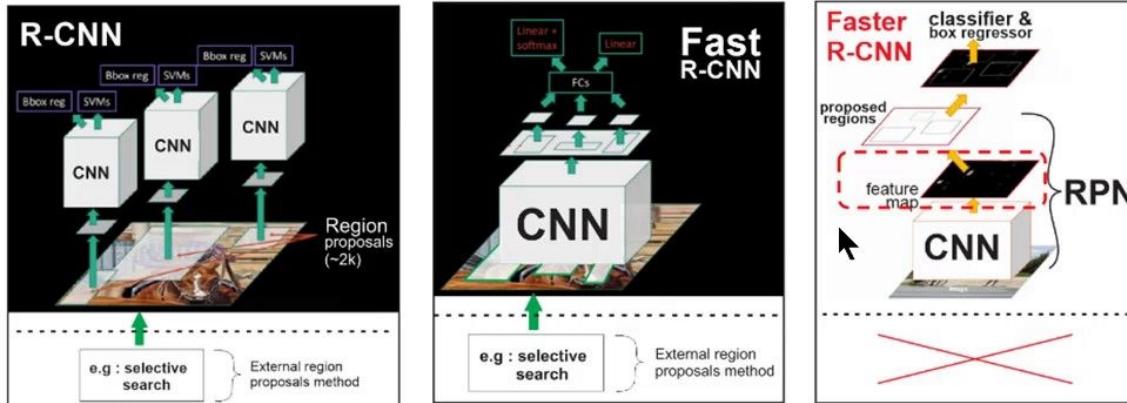
1. RPN classify object / not object
2. RPN regress box coordinates
3. Final classification score (object classes)
4. Final box coordinates

**Only one CNN  
needs to be  
trained** and we get  
region proposals  
almost for free



Ren et al., "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks", NIPS 2015  
Figure copyright 2015, Ross Girshick; reproduced with permission

## R-CNN variants



	R-CNN	Fast R-CNN	Faster R-CNN
Test time per image	50 seconds	2 seconds	0.2 seconds
Speed-up	1x	25x	250x
mAP (VOC 2007)	66.0%	66.9%	66.9%

\* Standford lecture notes on CNN by Fei Fei Li and Andrej Karpathy

5

By Ardiyan Umam



Demo time! Evaluating a pre trained COCO detector

Open in Colaboratory:

[https://colab.research.google.com/github/rnditdev/PyData\\_London\\_2018\\_Computer\\_Vision/blob/master/4A\\_Object\\_detection.ipynb](https://colab.research.google.com/github/rnditdev/PyData_London_2018_Computer_Vision/blob/master/4A_Object_detection.ipynb)



➤ Training your own  
Pizza detector with  
Faster-RCNN

## > Training setup



images



training



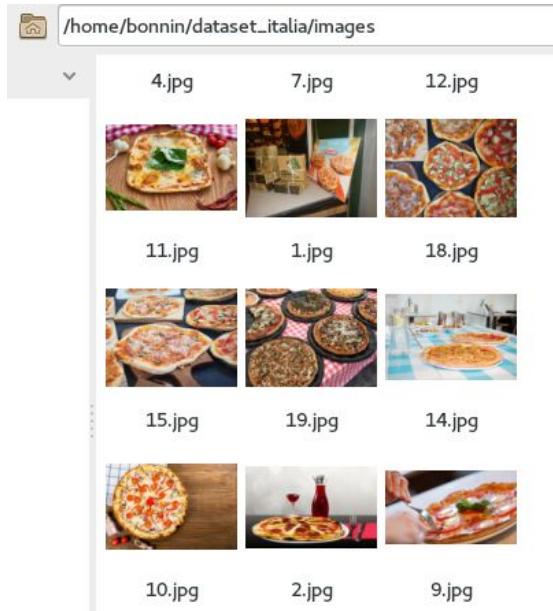
data



annotations



imgs\_url.txt



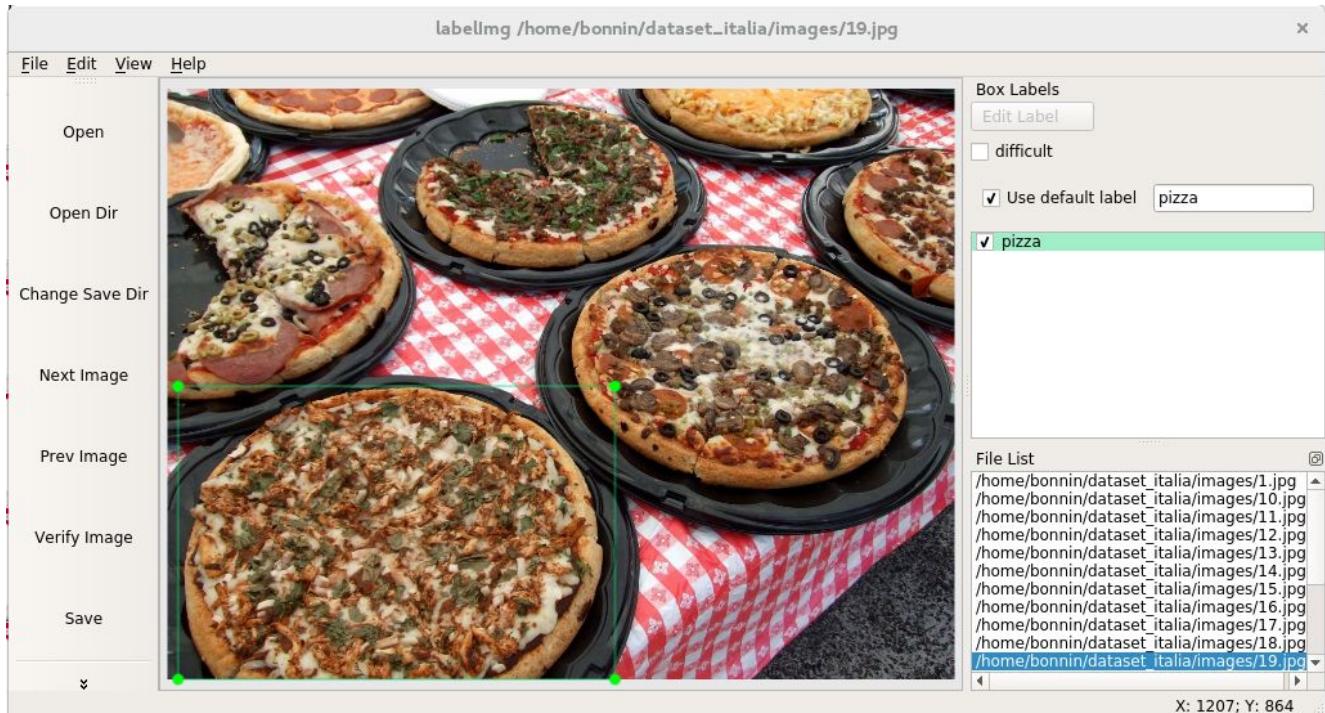
Training:  
object\_detection.pbtxt

```
item {  
  id: 1  
  name:  
  'pizza'  
}
```

Data:  
Empty(before  
training)

## ➤ Building the dataset: Labeling the images

```
sudo apt-get install pyqt5-dev-tools  
sudo pip3 install lxml  
make qt5py3  
python3 labelImg.py
```



## Model Training and final object snapshot saving

```
2 sudo apt-get install -y python-pip
3 sudo apt-get install -y protobuf-compiler python-pil python-lxml python-tk unzip
4 sudo pip install Cython
5 sudo pip install jupyter
6 sudo pip install matplotlib
7 sudo pip install pandas
8 sudo pip install tensorflow
9
10 git clone https://github.com/tensorflow/models.git
11 cd models
12 git checkout 079d67d9a0b3407e8d074a200780f3835413ef99
13 cd research
14 export PYTHONPATH=$PYTHONPATH:`pwd`:`pwd`/slim
15 protoc object_detection/protos/*.proto --python_out=.
16 python setup.py build
17 python setup.py install
18 cd object_detection
19 wget https://s3.amazonaws.com/italia18/dataset_italia_final.zip
20 unzip dataset_italia_final.zip
21
22 wget http://download.tensorflow.org/models/object_detection/faster_rcnn_inception_v2_coco_2018_01_28.tar.gz
23 tar -zxvf faster_rcnn_inception_v2_coco_2018_01_28.tar.gz
24
25 python xml_to_csv.py
26
27 python generate_tfrecord.py --csv_input=images/train_labels.csv --image_dir=images/train --output_path=train.record
28 python generate_tfrecord.py --csv_input=images/test_labels.csv --image_dir=images/test --output_path=test.record
29
30 python train.py --logtostderr --train_dir=training/ --pipeline_config_path=training/faster_rcnn_inception_v2_pets.config
31
32
```

## ➤ Demo Time: Inference of our trained model on test images

[https://colab.research.google.com/github/rnditdev/PyData\\_London\\_2018\\_Computer\\_Vision/blob/master/Object\\_detection\\_trained\\_model.ipynb](https://colab.research.google.com/github/rnditdev/PyData_London_2018_Computer_Vision/blob/master/Object_detection_trained_model.ipynb)

### Pizze d'Autore del Viandante

#### Mari e Monti:

Olio extravergine di Oliva, polpa di pomodoro, misto bosco, Gamberi dell'Adriatico, prezzemolo € 12,00



#### Profumi di bosco:

Mozzarella fiordilatte, scamorza affumicata, finferli e porcini freschi, fuori cottura: Bufala Campana e prezzemolo € 12,00



#### Rosetta:

Mozzarella fiordilatte, fuori cottura: Mortadella tartufata della Lessinia, Bufala Campana e mousse di pistacchi tostati € 12,00



#### Ortolana:

Melanze e zucchine grigliate, cipolla rossa di Tropea, pomodorini confit e Monte Veronese a scaglie, Fuori cottura: Bufala Campana e trito di erba cipollina € 12,00



#### Lavaredo:

Bufala Campana, crema di porcini, Speck tagliato a julienne, Grana Trentino a scaglie e un filo di miele € 11,00



#### Colonnata:

Mozzarella fiordilatte, Gorgonzola, Fuori cottura: radicchio julienné, Lardo di Colonnata, nodi sgusciate. € 12,00



#### Leggera:

Olio extravergine di Oliva, Polpa di pomodoro, olive taggiasche, cipolla rossa di Tropea, Fuori cottura: pomodorini confit, Bufala Campana, basilico julienné € 11,00



### Family # 2

\$28.99

1 extra-large two toppings

8 pcs Cajun chicken tenders

Large pizza: 99% wedges

Two liter soda



 **Image Segmentation**

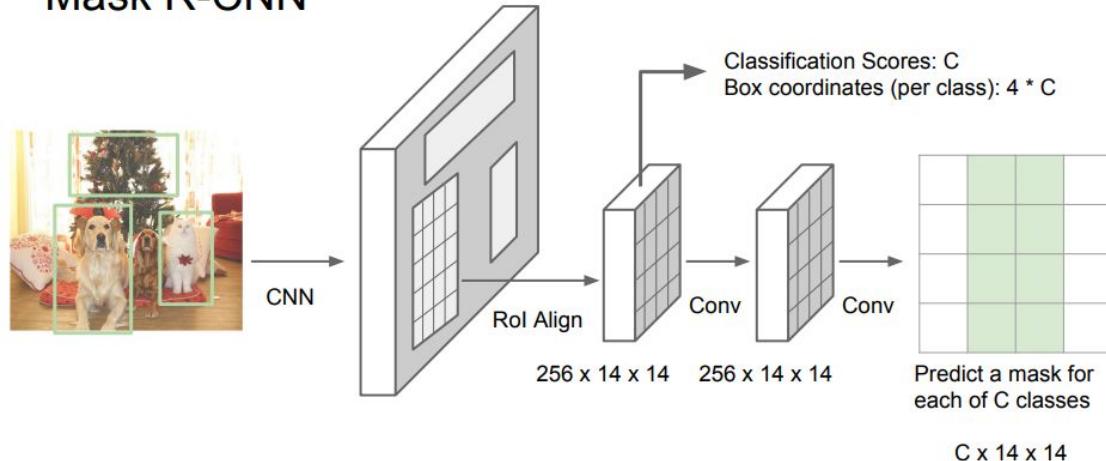
Adds a branch to Faster R-CNN that outputs a binary mask that says whether or not a given pixel is part of an object.

It's a Fully Convolutional Network on top of a CNN based feature map.

**Inputs:** CNN Feature Map.

**Outputs:** Matrix with 1s on all locations where the pixel belongs to the object and 0s elsewhere (this is known as a binary mask).

## Mask R-CNN





It consists of complex everyday scenes containing common objects in their natural context. Objects are labeled using per-instance segmentations to aid in precise object localization. Its canonical dataset has :

- **330K images (>200K labeled)**
- **1.5 million object instances**
- **80 object categories**
- **91 stuff categories**
- **5 captions per image**
- **250,000 people with keypoints**

The most common tasks applied to this dataset are:

- **Object segmentation ✓**
- **Recognition in context**
- **Superpixel stuff segmentation**

## ➤ MSCOCO general format

Coco json sample format:

```
{  
  "info" : info, "images" : [image], "annotations" : [annotation],  
  "licenses" : [license],  
}  
  
info{  
  "year" : int, "version" : str, "description" : str, "contributor" : str,  
  "url" : str, "date_created" : datetime,  
}  
  
image{  
  "id" : int, "width" : int, "height" : int, "file_name" : str, "license"  
  : int, "flickr_url" : str, "coco_url" : str, "date_captured" : datetime,  
}  
  
license{  
  "id" : int, "name" : str, "url" : str,  
}
```

## ➤ MSCOCO Annotation format

### Coco json annotation sample:

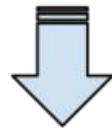
```
annotation{  
    "id" : int, "image_id" : int, "category_id" : int, "segmentation" : RLE  
    or [polygon], "area" : float, "bbox" : [x,y,width,height], "iscrowd" : 0  
    or 1,  
}  
  
categories[ {  
    "id" : int, "name" : str, "supercategory" : str,  
} ]
```

## ➤ MSCOCO Annotation format

### RLE



run-length encoding



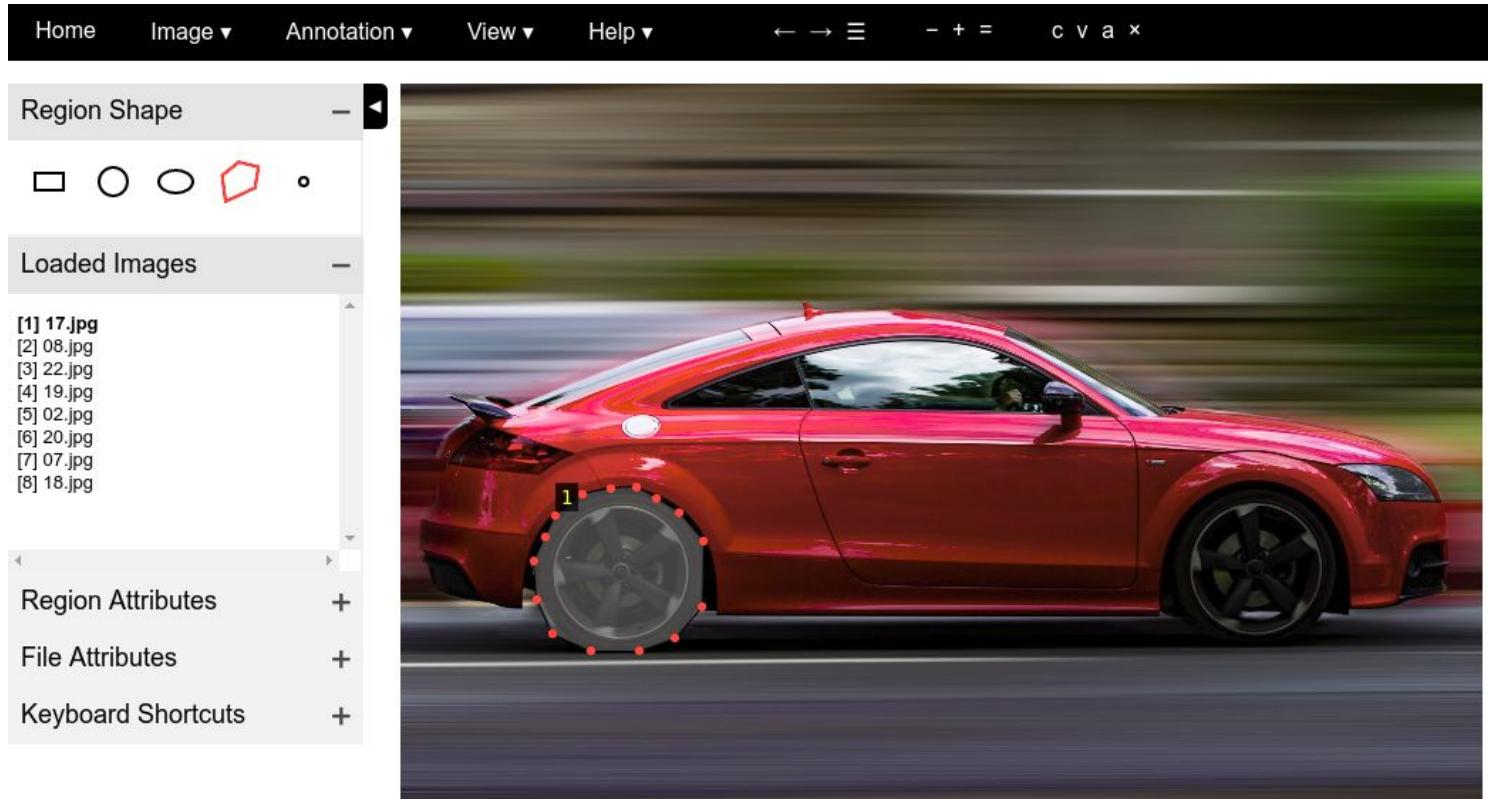
## ➤ Demo time: Evaluating a pre-trained COCO model

[https://colab.research.google.com/github/rnditdev/PyData\\_London\\_2018\\_Computer\\_Vision/  
blob/master/5A\\_Object Segmentation.ipynb](https://colab.research.google.com/github/rnditdev/PyData_London_2018_Computer_Vision/blob/master/5A_Object%20Segmentation.ipynb)



➤ Training a new  
Segmentation  
model.

# ➤ Mask-RCNN: Training a custom model: Tire segmenter



## ➤ Mask-RCNN: Training the model

[https://colab.research.google.com/github/rnditdev/PyData\\_London\\_2018\\_Computer\\_Vision/blob/master/5B\\_Object\\_segmentation\\_training.ipynb](https://colab.research.google.com/github/rnditdev/PyData_London_2018_Computer_Vision/blob/master/5B_Object_segmentation_training.ipynb)

## ➤ Segmenting test images

[https://colab.research.google.com/github/rnditdev/PyData\\_London\\_2018\\_Computer\\_Vision/  
blob/master/5B\\_Object\\_segmentation\\_training.ipynb](https://colab.research.google.com/github/rnditdev/PyData_London_2018_Computer_Vision/blob/master/5B_Object_segmentation_training.ipynb)





The future

Questions?



**machinalis**

Machine Learning Solutions Delivery